

# Comprehensive Analysis

```
import torch.nn as nn
from torchvision import datasets, transforms
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, classification_report
import os
import re
import time
import seaborn as sns

from src.custom_nn import CustomNeuralNet, train_model, setup_device
from src.logistic_regression import LogisticRegression
from src.multiclass_regression import SoftmaxRegression
from src.eval import *
```

## C1. Hyperparameter Analysis

### 1 Prepare MNIST dataset

```
train_data = datasets.MNIST(root="data", train=True,
transform=transforms.ToTensor(), download=True)
test_data = datasets.MNIST(root="data", train=False,
transform=transforms.ToTensor(), download=True)

mean = train_data.data.float().mean() / 255.0
std = train_data.data.float().std() / 255.0
print(f"Train data mean: {mean:.4f}, std: {std:.4f}")

X_train = (train_data.data.view(-1, 28*28).float() / 255.0 - mean) /
std
y_train = train_data.targets

X_test = (test_data.data.view(-1, 28*28).float() / 255.0 - mean) / std
y_test = test_data.targets

print("Train data shape:", X_train.shape)
print("Train labels shape:", y_train.shape)
print("Test data shape:", X_test.shape)
print("Test labels shape:", y_test.shape)

Train data mean: 0.1307, std: 0.3081
Train data shape: torch.Size([60000, 784])
Train labels shape: torch.Size([60000])
```

```
Test data shape: torch.Size([10000, 784])
Test labels shape: torch.Size([10000])
```

## 2 Learning Rate Analysis

Train different models

```
learning_rates = [0.001, 0.01, 0.1, 1.0]
batch_size = 64
results_dir = os.path.join("hyperparam_results", "lr")

for lr in learning_rates:
    print(f"\n==== Training with learning rate: {lr} ====")

    model = CustomNeuralNet(
        sizes=[28*28, 128, 64, 10],
        activation=nn.ReLU,
        weight_init="he"
    )

    lr_dir = os.path.join(results_dir, f"lr_{lr}")

    train_model(
        model,
        X=X_train,
        y=y_train,
        epochs=30,
        batch=batch_size,
        lr=lr,
        val_fraction=0.3,
        patience=5,
        destination=lr_dir
    )
```

Read CSVs into dictionary

```
learning_rates = [0.001, 0.01, 0.1, 1.0]
batch_size = 64
results_dir = os.path.join("hyperparam_results", "lr")
results = {}
for lr in learning_rates:
    csv_path = os.path.join(results_dir, f"lr_{lr}", "model1",
"results.csv")
    if os.path.exists(csv_path):
        results[lr] = pd.read_csv(csv_path)
    else:
        print(f"Warning: Missing {csv_path}")
```

Figure 1: Learning Curves (Loss & Accuracy)

```

fig, axes = plt.subplots(2, len(results), figsize=(18, 8))
plt.suptitle("Learning Rate Analysis – Loss and Accuracy per LR",
font-size=16, font-weight="bold")

for i, (lr, df) in enumerate(results.items()):
    # Loss
    ax_loss = axes[0, i]
    ax_loss.plot(df["epoch"], df["train_loss"], label="Train Loss",
color="blue")
    ax_loss.plot(df["epoch"], df["val_loss"], '--', label="Val Loss",
color="orange")
    ax_loss.set_title(f"LR = {lr}")
    ax_loss.set_xlabel("Epoch")
    ax_loss.set_ylabel("Loss")
    ax_loss.grid(True, linestyle="--", alpha=0.5)
    if i == 0:
        ax_loss.legend(fontsize=8)

    # Accuracy
    ax_acc = axes[1, i]
    ax_acc.plot(df["epoch"], df["train_acc"], label="Train Acc",
color="green")
    ax_acc.plot(df["epoch"], df["val_acc"], '--', label="Val Acc",
color="red")
    ax_acc.set_xlabel("Epoch")
    ax_acc.set_ylabel("Accuracy")
    ax_acc.grid(True, linestyle="--", alpha=0.5)
    if i == 0:
        ax_acc.legend(fontsize=8)

plt.tight_layout(rect=(0, 0, 1, 0.95))
plt.show()

```





Figure 2: Convergence Analysis (Speed & Stability)

```
plt.figure(figsize=(8, 6))
plt.title("Convergence Speed and Stability", fontsize=14,
fontweight="bold")

for lr, df in results.items():
    # Compute per-epoch change in val_loss to estimate convergence
    # rate
    val_loss_change = df["val_loss"].diff().abs()
    plt.plot(df["epoch"], val_loss_change, label=f"lr={lr}")

plt.xlabel("Epoch")
plt.ylabel("|\Delta Validation Loss| (Change per Epoch)")
plt.legend()
plt.grid(True)
plt.show()
```



## Analysis and Comparing Models

```
summary = []
for lr, df in results.items():
    best_idx = df["val_acc"].idxmax()
    best_row = df.iloc[best_idx]

    summary.append({
        "Learning Rate": lr,
        "Best Epoch": best_idx + 1,
        "Best Train Loss": best_row["train_loss"],
        "Best Val Loss": best_row["val_loss"],
        "Best Train Acc": best_row["train_acc"],
        "Best Val Acc": best_row["val_acc"],
    })

summary_df = pd.DataFrame(summary)
summary_df = summary_df.sort_values(by="Best Val Acc",
                                    ascending=False).reset_index(drop=True)

# Highlight best row
best_row = summary_df.iloc[0]
best_lr = best_row["Learning Rate"]

print("==== Learning Rate Performance Summary ===")
display(
    summary_df.style
        .highlight_max(subset=["Best Val Acc"], color="green")
        .format(precision=4)
        .set_table_styles([{'selector': 'th', 'props': [('font-weight', 'bold')] }])
        .hide(axis="index")
)
==== Learning Rate Performance Summary ===
<pandas.io.formats.style.Styler at 0x27e821e7c70>
```

## 2 Batch Size Analysis

Train different models

```
batch_sizes = [8, 16, 32, 64, 128, 256]
results_dir = os.path.join("hyperparam_results", "batch_size")

for batch in batch_sizes:
    print(f"\n==== Training with batch size: {batch} ===")

    model = CustomNeuralNet(
        sizes=[28*28, 128, 64, 10],
        activation=nn.ReLU,
```

```

        weight_init="he"
    )

batch_dir = os.path.join(results_dir, f"batch_{batch}")

train_model(
    model,
    X=X_train,
    y=y_train,
    epochs=30,
    batch=batch,
    lr=0.01,
    val_fraction=0.3,
    patience=5,
    destination=batch_dir
)

```

### Read CSVs into dictionary

```

batch_sizes = [8, 16, 32, 64, 128, 256]
results_dir = os.path.join("hyperparam_results", "batch_size")
results = {}
for bs in batch_sizes:
    path = os.path.join(results_dir, f"batch_{bs}", "model1",
"results.csv")
    if os.path.exists(path):
        results[bs] = pd.read_csv(path)
    else:
        print(f"Warning: Missing {path}")

```

### Plot learning curves for all batch sizes

```

n = len(results)
fig, axes = plt.subplots(2, n, figsize=(4*n, 8))
plt.suptitle("Batch Size Analysis – Learning Curves", fontsize=14,
fontweight="bold")

for i, (bs, df) in enumerate(results.items()):
    # Loss
    axes[0][i].plot(df["epoch"], df["train_loss"], label="Train Loss",
color="tab:blue")
    axes[0][i].plot(df["epoch"], df["val_loss"], label="Val Loss",
color="tab:orange", linestyle="--")
    axes[0][i].set_title(f"Batch Size = {bs}")
    axes[0][i].set_xlabel("Epoch")
    axes[0][i].set_ylabel("Loss")
    axes[0][i].grid(True)
    axes[0][i].legend()

    # Accuracy


```

```

        axes[1][i].plot(df["epoch"], df["train_acc"], label="Train Acc",
color="tab:green")
        axes[1][i].plot(df["epoch"], df["val_acc"], label="Val Acc",
color="tab:red", linestyle="--")
        axes[1][i].set_xlabel("Epoch")
        axes[1][i].set_ylabel("Accuracy")
        axes[1][i].grid(True)
        axes[1][i].legend()

plt.tight_layout(rect=(0, 0, 1, 0.95))
plt.show()

```



## Convergence analysis and comparing models

```

summary = []
for bs, df in results.items():
    best_idx = df["val_acc"].idxmax()
    best_row = df.iloc[best_idx]

    summary.append({
        "Batch Size": bs,
        "Best Epoch": best_idx + 1,
        "Best Train Loss": best_row["train_loss"],
        "Best Val Loss": best_row["val_loss"],
        "Best Train Acc": best_row["train_acc"],
        "Best Val Acc": best_row["val_acc"],
    })

summary_df = pd.DataFrame(summary)
summary_df = summary_df.sort_values(by="Best Val Acc",
ascending=False).reset_index(drop=True)

print("==== Batch Size Performance Summary ====")
display(
    summary_df.style
        .highlight_max(subset=[ "Best Val Acc"], color="green")
        .format(precision=4)
)

```

```

        .set_table_styles([{'selector': 'th', 'props': [('font-weight', 'bold')]})
        .hide(axis="index")
    )

==== Batch Size Performance Summary ====
<pandas.io.formats.style.Styler at 0x27e82ed64d0>

```

## 3 Architecture Analysis

Train different models

```

hidden_layers_list = [1, 2, 3, 4, 5]
neurons_list = [64, 128, 256, 512]
results_dir = os.path.join("hyperparam_results", "arch")

for n_layers in hidden_layers_list:
    for n_neurons in neurons_list:
        sizes = [28*28] + [n_neurons]*n_layers + [10]
        print(f"\n==== Architecture: {sizes} ====")

        model = CustomNeuralNet(
            sizes=sizes,
            activation=nn.ReLU,
            weight_init="he"
        )

        arch_dir = os.path.join(results_dir,
f"arch_{n_layers}L_{n_neurons}N")
        train_model(
            model,
            X=X_train,
            y=y_train,
            epochs=30,
            batch=64,
            lr=0.01,
            val_fraction=0.3,
            patience=5,
            destination=arch_dir
        )
    
```

Read CSVs into dictionary

```

def extract_numbers(folder_name):
    nums = re.findall(r'\d+', folder_name)
    return tuple(map(int, nums)) if nums else (0,)

hidden_layers_list = [1, 2, 3, 4, 5]

```

```

neurons_list = [64, 128, 256, 512]
results_dir = os.path.join("hyperparam_results", "arch")
results = {}
for folder in sorted(os.listdir(results_dir), key=lambda f:
extract_numbers(f)):
    path = os.path.join(results_dir, folder, "modell", "results.csv")
    if os.path.exists(path):
        results[folder] = pd.read_csv(path)
    else:
        print(f"Warning: Missing {path}")

```

Plot learning curves for all architectures

```

n = len(results)
fig, axes = plt.subplots(n, 2, figsize=(12, 3*n))
plt.suptitle("Architecture Analysis – Learning Curves", fontsize=14,
fontweight="bold")

if n == 1:
    axes = [axes]


for i, (arch, df) in enumerate(results.items()):
    # --- Loss ---
    axes[i][0].plot(df["epoch"], df["train_loss"], label="Train Loss",
color="tab:blue")
    axes[i][0].plot(df["epoch"], df["val_loss"], label="Val Loss",
color="tab:orange", linestyle="--")
    axes[i][0].set_title(f"{arch}: Loss")
    axes[i][0].set_xlabel("Epoch")
    axes[i][0].set_ylabel("Loss")
    axes[i][0].grid(True)
    axes[i][0].legend()

    # --- Accuracy ---
    axes[i][1].plot(df["epoch"], df["train_acc"], label="Train Acc",
color="tab:green")
    axes[i][1].plot(df["epoch"], df["val_acc"], label="Val Acc",
color="tab:red", linestyle="--")
    axes[i][1].set_title(f"{arch}: Accuracy")
    axes[i][1].set_xlabel("Epoch")
    axes[i][1].set_ylabel("Accuracy")
    axes[i][1].grid(True)
    axes[i][1].legend()

plt.tight_layout(rect=(0, 0, 1, 0.95))
plt.show()

```

## Architecture Analysis — Learning Curves



## Analysis and Comparing Models

```
summary = []
for arch, df in results.items():
    best_idx = df["val_acc"].idxmax()
    best_row = df.iloc[best_idx]

    summary.append({
        "Architecture": arch,
        "Best Epoch": int(best_row["epoch"]),
        "Best Train Acc": best_row["train_acc"],
        "Best Val Acc": best_row["val_acc"],
        "Train Loss": best_row["train_loss"],
        "Val Loss": best_row["val_loss"]
    })

summary_df = pd.DataFrame(summary).reset_index(drop=True)
best_idx = summary_df["Best Val Acc"].idxmax()

def highlight_best_row(row):
    color = 'blue' if row.name == best_idx else ''
    return ['background-color: {}'.format(color)] * len(row)

display(
    summary_df.style
        .apply(highlight_best_row, axis=1)
        .highlight_max(subset=["Best Val Acc"], color="green")
        .format(precision=4)
        .hide(axis="index")
)
<pandas.io.formats.style.Styler at 0x27e8205f190>
```

## 4 Train Best Neural Network

Based on analysis, best neural neural network with lr=0.1 batch\_size=8 arch=3 hidden layers with 512 neuron

```
sizes = [28*28] + [513]*3 + [10]
lr=0.1
batch_sz=8

model = CustomNeuralNet(
    sizes=sizes,
    activation=nn.ReLU,
    weight_init="he"
)

dirct = os.path.join("classification_models", "nn")
```

```

train_model(
    model,
    X=X_train,
    y=y_train,
    epochs=30,
    batch=batch_sz,
    lr=lr,
    val_fraction=0.3,
    patience=5,
    destination=dirct
)

```

## C2. Model Comparison

Compare logistic regression, softmax regression, and best neural network

### 1 Prepare MNIST test dataset

```

train_data = datasets.MNIST(root="data", train=True,
transform=transforms.ToTensor(), download=True)
test_data = datasets.MNIST(root="data", train=False,
transform=transforms.ToTensor(), download=True)

mean = train_data.data.float().mean() / 255.0
std = train_data.data.float().std() / 255.0

X_test = (test_data.data.view(-1, 28*28).float() / 255.0 - mean) / std
y_test = test_data.targets

# Binary model: 0 vs 1 classification
binary_indices = (y_test == 0) | (y_test == 1)
X_test_binary = X_test[binary_indices]
y_test_binary = y_test[binary_indices]

print("Test data shape:", X_test.shape)
print("Test labels shape:", y_test.shape)
print(f"\nBinary test set size: {len(y_test_binary)} samples")
print(f"Multiclass test set size: {len(y_test)} samples")

Test data shape: torch.Size([10000, 784])
Test labels shape: torch.Size([10000])

Binary test set size: 2115 samples
Multiclass test set size: 10000 samples

```

## 2 Load Models

```
bin_model_path = os.path.join("classification_models",
"binary_logistic_regression_model.pth")
multi_model_path = os.path.join("classification_models",
"multi_class_softmax_regression_model.pth")
nn_model_path = os.path.join("classification_models", "nn", "model1",
"weights", "best.pt")

sizes = [28*28] + [513]*3 + [10]
device = setup_device()

# Load binary class model
try:
    binary_model = BinaryLogisticRegression()
    binary_checkpoint = torch.load(bin_model_path,
map_location=device)
    binary_model.load_state_dict(binary_checkpoint)
    binary_model.to(device)
    binary_model.eval()
    print("Binary model loaded successfully")
except Exception as e:
    print(f"Error loading binary model: {e}")
    binary_model = None

# Load multiclass model
try:
    multiclass_model = MulticlassLogisticRegression()
    multiclass_checkpoint = torch.load(multi_model_path,
map_location=device)
    multiclass_model.load_state_dict(multiclass_checkpoint)
    multiclass_model.to(device)
    multiclass_model.eval()
    print("Multiclass model loaded successfully")
except Exception as e:
    print(f"Error loading multiclass model: {e}")
    multiclass_model = None

# Load best neural network model
try:
    best_nn_model = CustomNeuralNet(
        sizes=sizes,
        activation=nn.ReLU,
        weight_init="he"
    )
    nn_checkpoint = torch.load(nn_model_path, map_location=device)
    best_nn_model.load_state_dict(nn_checkpoint['model_state'])
    best_nn_model.to(device)
    best_nn_model.eval()
    print("Neural Network model loaded successfully")
```

```
except Exception as e:
    print(f"Error loading NN model: {e}")
    best_nn_model = None

Utilizing device: cuda:0
Binary model loaded successfully
Multiclass model loaded successfully
Neural Network model loaded successfully

C:\Users\Youssef\AppData\Local\Temp\ipykernel_788\772945862.py:7:
FutureWarning: You are using `torch.load` with `weights_only=False`  

(the current default value), which uses the default pickle module  

implicitly. It is possible to construct malicious pickle data which  

will execute arbitrary code during unpickling (See  

https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models for more details). In a future release, the default value for  

`weights_only` will be flipped to `True`. This limits the functions  

that could be executed during unpickling. Arbitrary objects will no  

longer be allowed to be loaded via this mode unless they are  

explicitly allowlisted by the user via  

`torch.serialization.add_safe_globals`. We recommend you start setting  

`weights_only=True` for any use case where you don't have full control  

of the loaded file. Please open an issue on GitHub for any issues  

related to this experimental feature.

    binary_checkpoint = torch.load(bin_model_path, map_location=device)
C:\Users\Youssef\AppData\Local\Temp\ipykernel_788\772945862.py:19:
FutureWarning: You are using `torch.load` with `weights_only=False`  

(the current default value), which uses the default pickle module  

implicitly. It is possible to construct malicious pickle data which  

will execute arbitrary code during unpickling (See  

https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models for more details). In a future release, the default value for  

`weights_only` will be flipped to `True`. This limits the functions  

that could be executed during unpickling. Arbitrary objects will no  

longer be allowed to be loaded via this mode unless they are  

explicitly allowlisted by the user via  

`torch.serialization.add_safe_globals`. We recommend you start setting  

`weights_only=True` for any use case where you don't have full control  

of the loaded file. Please open an issue on GitHub for any issues  

related to this experimental feature.

    multiclass_checkpoint = torch.load(multi_model_path,
map_location=device)
C:\Users\Youssef\AppData\Local\Temp\ipykernel_788\772945862.py:35:
FutureWarning: You are using `torch.load` with `weights_only=False`  

(the current default value), which uses the default pickle module  

implicitly. It is possible to construct malicious pickle data which  

will execute arbitrary code during unpickling (See  

https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models for more details). In a future release, the default value for  

`weights_only` will be flipped to `True`. This limits the functions
```

```
that could be executed during unpickling. Arbitrary objects will no longer be allowed to be loaded via this mode unless they are explicitly allowlisted by the user via `torch.serialization.add_safe_globals`. We recommend you start setting `weights_only=True` for any use case where you don't have full control of the loaded file. Please open an issue on GitHub for any issues related to this experimental feature.  
nn_checkpoint = torch.load(nn_model_path, map_location=device)
```

## 3 Evaluate Models

Run evaluation

```
results_comparison = {}
```

Binary Model

```
start_time = time.time()  
acc, loss, preds, labels = evaluate_model(  
    binary_model, X_test_binary, y_test_binary, device,  
    model_type='binary'  
)  
eval_time = time.time() - start_time  
  
num_params = sum(p.numel() for p in binary_model.parameters())  
  
results_comparison['Binary Logistic Regression'] = {  
    'accuracy': acc,  
    'loss': loss,  
    'num_params': num_params,  
    'eval_time': eval_time,  
    'predictions': preds,  
    'labels': labels,  
    'test_size': len(y_test_binary),  
    'task': 'Binary (0 vs 1)'  
}
```

Multi-Class Model

```
start_time = time.time()  
acc, loss, preds, labels = evaluate_model(  
    multiclass_model, X_test, y_test, device, model_type='multiclass'  
)  
eval_time = time.time() - start_time  
  
num_params = sum(p.numel() for p in multiclass_model.parameters())  
  
results_comparison['Softmax Regression'] = {  
    'accuracy': acc,
```

```

    'loss': loss,
    'num_params': num_params,
    'eval_time': eval_time,
    'predictions': preds,
    'labels': labels,
    'test_size': len(y_test),
    'task': 'Multiclass (10 digits)'
}

```

## Neural-Network Model

```

start_time = time.time()
acc, loss, preds, labels = evaluate_model(
    best_nn_model, X_test, y_test, device, model_type='multiclass'
)
eval_time = time.time() - start_time

num_params = sum(p.numel() for p in best_nn_model.parameters())

results_comparison['Best Neural Network'] = {
    'accuracy': acc,
    'loss': loss,
    'num_params': num_params,
    'eval_time': eval_time,
    'predictions': preds,
    'labels': labels,
    'test_size': len(y_test),
    'task': 'Multiclass (10 digits)'
}

```

## Performance Summary Table

```

print("\n" + "*80")
print("COMPREHENSIVE PERFORMANCE SUMMARY TABLE")
print("*80")

summary_data = []
for name, data in results_comparison.items():
    # Calculate computational complexity (FLOPs approximation)
    if 'Binary' in name:
        flops = 784 * 1 * 2 # Input * Output * (mult + add)
        complexity = "Very Low"
    elif 'Softmax' in name:
        flops = 784 * 10 * 2
        complexity = "Low"
    else:
        flops = 0
        prev = 784
        for i in range(1, len(data.get('architecture', []))-1):
            flops += prev * data['architecture'][i]
            prev = data['architecture'][i]
        complexity = "High"
    summary_data.append({
        'name': name,
        'flops': flops,
        'complexity': complexity
    })

```

## Performance Visualization

```
fig, axes = plt.subplots(1, 3, figsize=(18, 5))
fig.suptitle('Model Comparison - Performance Metrics', fontsize=16,
```

```

fontweight='bold')

model_names = list(results_comparison.keys())
# Only multiclass models
multiclass_names = [name for name in model_names if 'Binary' not in name]

if len(multiclass_names) >= 2:
    multiclass_accs = [results_comparison[m]['accuracy'] for m in multiclass_names]
    multiclass_params = [results_comparison[m]['num_params'] for m in multiclass_names]
    multiclass_times = [results_comparison[m]['eval_time'] for m in multiclass_names]

# Plot 1: Accuracy Comparison
colors = ['#3498db', '#e74c3c', '#2ecc71']
bars = axes[0].bar(range(len(multiclass_names)), multiclass_accs,
color=colors[:len(multiclass_names)], alpha=0.7)
axes[0].set_xlabel('Model', fontweight='bold')
axes[0].set_ylabel('Test Accuracy (%)', fontweight='bold')
axes[0].set_title('Test Accuracy Comparison', fontweight='bold')
axes[0].set_xticks(range(len(multiclass_names)))
axes[0].set_xticklabels([name.replace(' ', '\n') for name in multiclass_names], fontsize=9)
axes[0].grid(True, alpha=0.3, axis='y')

# Add value labels
for i, (bar, acc) in enumerate(zip(bars, multiclass_accs)):
    axes[0].text(bar.get_x() + bar.get_width()/2., acc + 0.5,
                 f'{acc:.2f}%', ha='center', va='bottom',
                 fontweight='bold')

# Plot 2: Accuracy vs Parameters
axes[1].scatter(multiclass_params, multiclass_accs, s=200,
alpha=0.6,
                 c=range(len(multiclass_names)), cmap='viridis')

for i, name in enumerate(multiclass_names):
    axes[1].annotate(name.replace(' ', '\n'),
                     (multiclass_params[i], multiclass_accs[i]),
                     fontsize=8, ha='center', xytext=(0, 10),
                     textcoords='offset points')

axes[1].set_xlabel('Number of Parameters', fontweight='bold')
axes[1].set_ylabel('Test Accuracy (%)', fontweight='bold')
axes[1].set_title('Accuracy vs Model Complexity',
fontweight='bold')
axes[1].grid(True, alpha=0.3)

```


```

# Plot 3: Speed Comparison (Samples per second)
speeds = [results_comparison[m]['test_size']/results_comparison[m]
['eval_time']]
for m in multiclass_names]
bars = axes[2].barh(range(len(multiclass_names)), speeds,
color=colors[:len(multiclass_names)],
alpha=0.7)
axes[2].set_ylabel('Model', fontweight='bold')
axes[2].set_xlabel('Inference Speed (samples/s)', fontweight='bold')
axes[2].set_title('Inference Speed Comparison', fontweight='bold')
axes[2].set_yticks(range(len(multiclass_names)))
axes[2].set_yticklabels(multiclass_names, fontsize=9)
axes[2].grid(True, alpha=0.3, axis='x')

# Add value labels
for i, (bar, speed) in enumerate(zip(bars, speeds)):
    axes[2].text(speed + 100, bar.get_y() + bar.get_height()/2.,
f'{speed:.0f}', va='center', fontweight='bold')

plt.tight_layout()
plt.show()

```



### Confusion Matrix for Best Model

```

best_model_name = max(
    [k for k in results_comparison.keys() if 'Binary' not in k],
    key=lambda x: results_comparison[x]['accuracy'])
)
best_data = results_comparison[best_model_name]

print(f"\n{'='*80}")
print(f"DETAILED EVALUATION: {best_model_name}")
print(f"{'='*80}")

# Confusion Matrix
cm = confusion_matrix(best_data['labels'], best_data['predictions'])

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 6))

```

```

# Absolute counts
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=ax1,
            xticklabels=range(10), yticklabels=range(10),
            cbar_kws={'label': 'Count'})
ax1.set_title(f'{best_model_name}\nConfusion Matrix - Absolute Counts',
              fontsize=13, fontweight='bold')
ax1.set_ylabel('True Label', fontsize=11, fontweight='bold')
ax1.set_xlabel('Predicted Label', fontsize=11, fontweight='bold')

# Normalized percentages
cm_normalized = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis] * 100
sns.heatmap(cm_normalized, annot=True, fmt='.1f', cmap='RdYlGn',
            ax=ax2,
            xticklabels=range(10), yticklabels=range(10),
            cbar_kws={'label': 'Percentage (%)'}, vmin=0, vmax=100)
ax2.set_title(f'{best_model_name}\nConfusion Matrix - Normalized (%)',
              fontsize=13, fontweight='bold')
ax2.set_ylabel('True Label', fontsize=11, fontweight='bold')
ax2.set_xlabel('Predicted Label', fontsize=11, fontweight='bold')

plt.tight_layout()
plt.show()

```

---




---

## DETAILED EVALUATION: Best Neural Network

---



---



## Classification Report

```
print("\n" + "-"*80)
print("DETAILED CLASSIFICATION REPORT")
print("-"*80)
print(classification_report(best_data['labels'],
best_data['predictions'],
digits=4, target_names=[str(i) for i in
range(10)]))
```

### DETAILED CLASSIFICATION REPORT

	precision	recall	f1-score	support
0	0.9836	0.9786	0.9811	980
1	0.9834	0.9903	0.9868	1135
2	0.9900	0.9641	0.9769	1032
3	0.9603	0.9822	0.9711	1010
4	0.9731	0.9562	0.9646	982
5	0.9693	0.9563	0.9628	892
6	0.9622	0.9833	0.9726	958
7	0.9553	0.9767	0.9658	1028
8	0.9627	0.9538	0.9582	974
9	0.9552	0.9504	0.9528	1009
accuracy			0.9696	10000
macro avg	0.9695	0.9692	0.9693	10000
weighted avg	0.9697	0.9696	0.9696	10000

## Misclassification Analysis

```
print("\n" + "="*80)
print("MISCLASSIFICATION ANALYSIS")
print("="*80)

preds = best_data['predictions']
labels = best_data['labels']

misclassified_indices = np.where(preds != labels)[0]
total_errors = len(misclassified_indices)

print(f"\nTotal Misclassified: {total_errors} out of {len(labels)}
({100*total_errors/len(labels):.2f}%)")

# Per-class error analysis
print("\n" + "-"*80)
```

```

print("PER-CLASS ERROR ANALYSIS")
print("-"*80)
print(f"{'Digit':<8} {'Total':<8} {'Correct':<8} {'Errors':<8}\n"
      {'Accuracy':<12} {'Error Rate':<12}"')
print("-"*80)

class_metrics = []
for digit in range(10):
    digit_mask = labels == digit
    digit_total = np.sum(digit_mask)
    digit_correct = np.sum((labels == digit) & (preds == digit))
    digit_errors = digit_total - digit_correct
    digit_acc = 100 * digit_correct / digit_total if digit_total > 0
    else 0
    error_rate = 100 * digit_errors / digit_total if digit_total > 0
    else 0

    print(f"{'digit':<8} {'digit_total':<8} {'digit_correct':<8}\n"
          {'digit_errors':<8} "
              f"{'digit_acc':<12.2f} {'error_rate':<12.2f}"')

    class_metrics.append({
        'digit': digit,
        'total': digit_total,
        'errors': digit_errors,
        'error_rate': error_rate
    })

# Most confused pairs
print("\n" + "-"*80)
print("MOST COMMON MISCLASSIFICATION PAIRS")
print("-"*80)

confusion_pairs = []
for i in range(10):
    for j in range(10):
        if i != j and cm[i, j] > 0:
            confusion_pairs.append((i, j, cm[i, j]))

confusion_pairs.sort(key=lambda x: x[2], reverse=True)

print(f"{'True Label':<12} {'Predicted As':<15} {'Count':<10} {'% of\n"
      'True Class':<15}"')
print("-"*80)
for true_label, pred_label, count in confusion_pairs[:15]:
    total_true = np.sum(labels == true_label)
    percentage = 100 * count / total_true
    print(f"{'true_label':<12} {'pred_label':<15} {'count':<10}\n"
          {'percentage:<15.2f}%")

```

---

---

---

## MISCLASSIFICATION ANALYSIS

---

---

Total Misclassified: 304 out of 10000 (3.04%)

---

---

## PER-CLASS ERROR ANALYSIS

---

---

Digit	Total	Correct	Errors	Accuracy	Error Rate
0	980	959	21	97.86	2.14
1	1135	1124	11	99.03	0.97
2	1032	995	37	96.41	3.59
3	1010	992	18	98.22	1.78
4	982	939	43	95.62	4.38
5	892	853	39	95.63	4.37
6	958	942	16	98.33	1.67
7	1028	1004	24	97.67	2.33
8	974	929	45	95.38	4.62
9	1009	959	50	95.04	4.96

---

---

## MOST COMMON MISCLASSIFICATION PAIRS

---

---

True Label	Predicted As	Count	% of True Class	%
4	9	27	2.75	%
5	6	13	1.46	%
2	7	12	1.16	%
5	3	12	1.35	%
9	3	12	1.19	%
0	7	10	1.02	%
8	6	10	1.03	%
2	8	9	0.87	%
9	7	9	0.89	%
9	4	8	0.79	%
9	5	8	0.79	%
8	3	7	0.72	%
8	4	7	0.72	%

2	1	6	0.58	%
4	6	6	0.61	%